A1 | ResNet-18 on CIFAR-100

Source: A1/run_a1.py + torchvision.models.resnet18 • 11,227,812 parameters

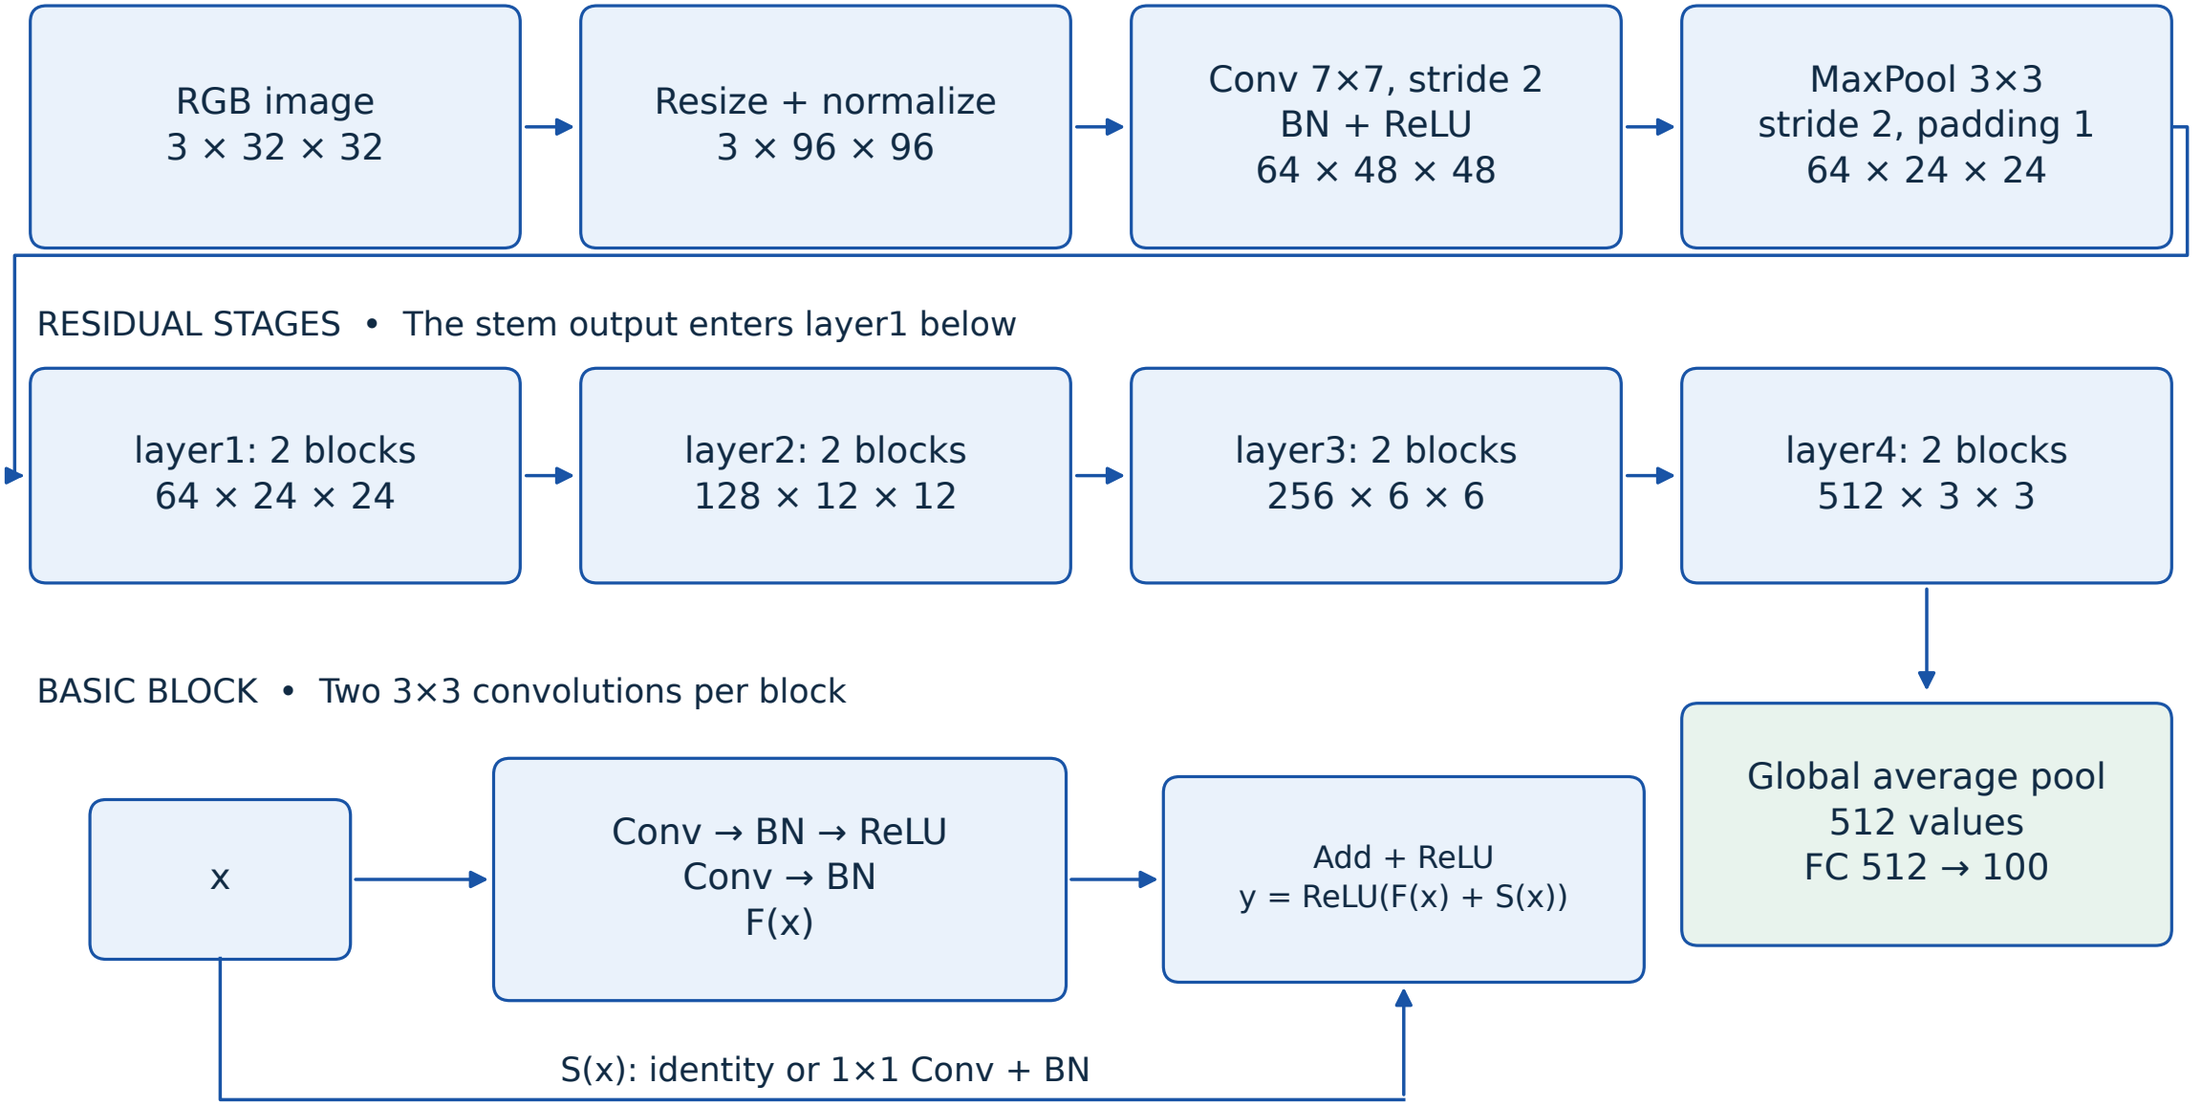

First blocks of layer2–4 downsample with stride 2 and a projected shortcut.
Head output: 100 logits. The original ImageNet stem is retained (not a CIFAR-specific stem).
Partial fine-tuning trains layer4 + fc; frozen BatchNorm statistics remain fixed.