E1 | Three image classifiers

Fashion-MNIST • Shapes omit the batch dimension • Source: E1/run_e1.py

SOFTMAX LINEAR • 7,850 parameters

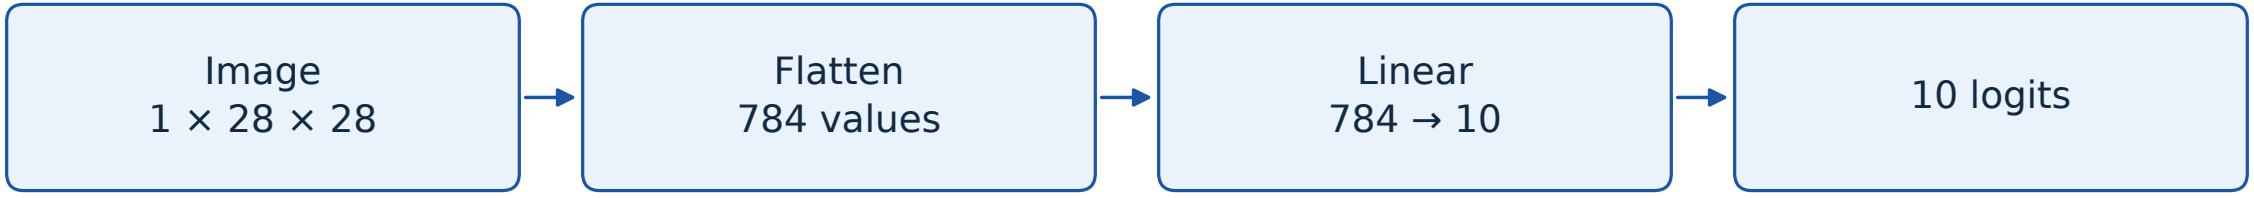

MLP • 109,386 parameters • Same image input

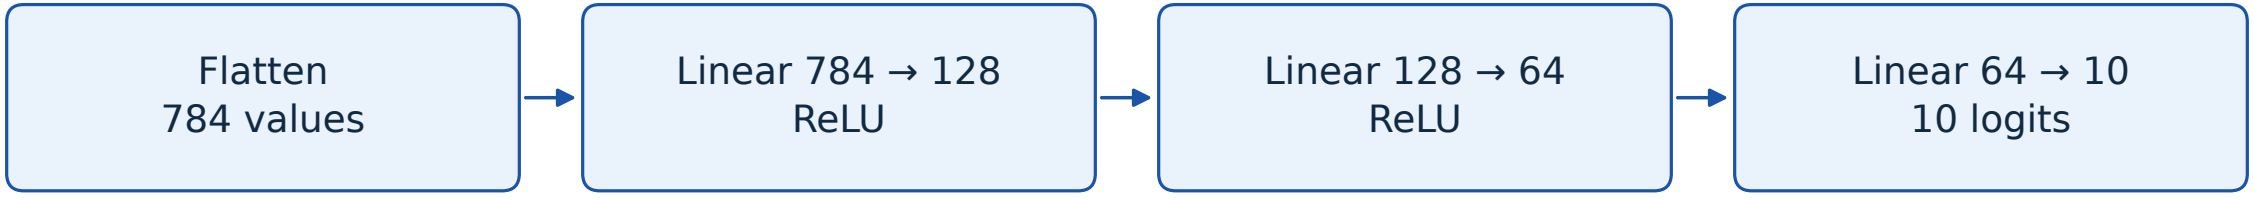

CNN • 105,866 parameters • Same image input

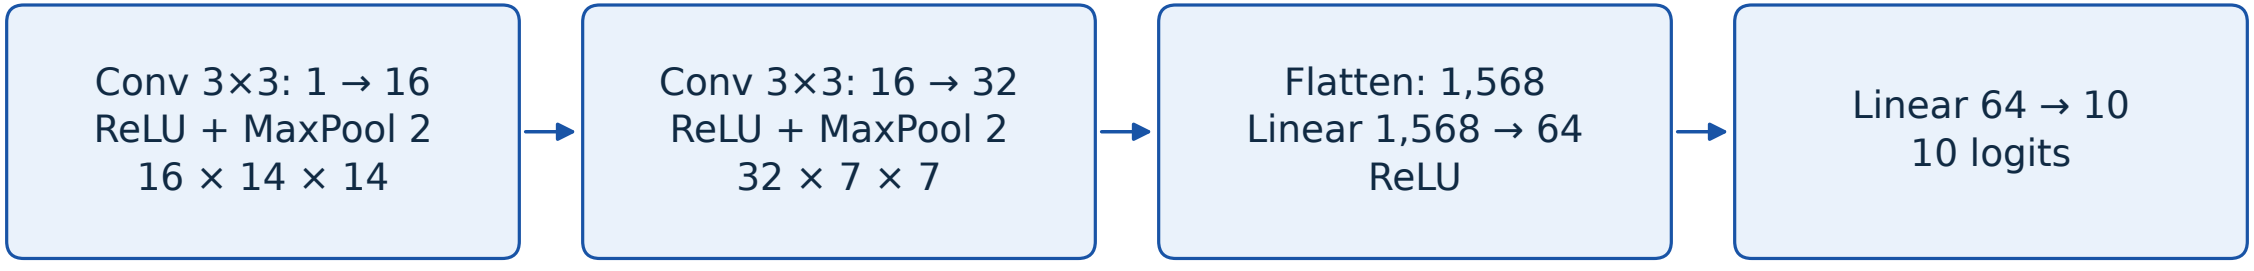

Convolutions: stride 1, padding 1. Pooling: kernel 2, stride 2.
All heads return raw logits; CrossEntropyLoss applies log-softmax internally.